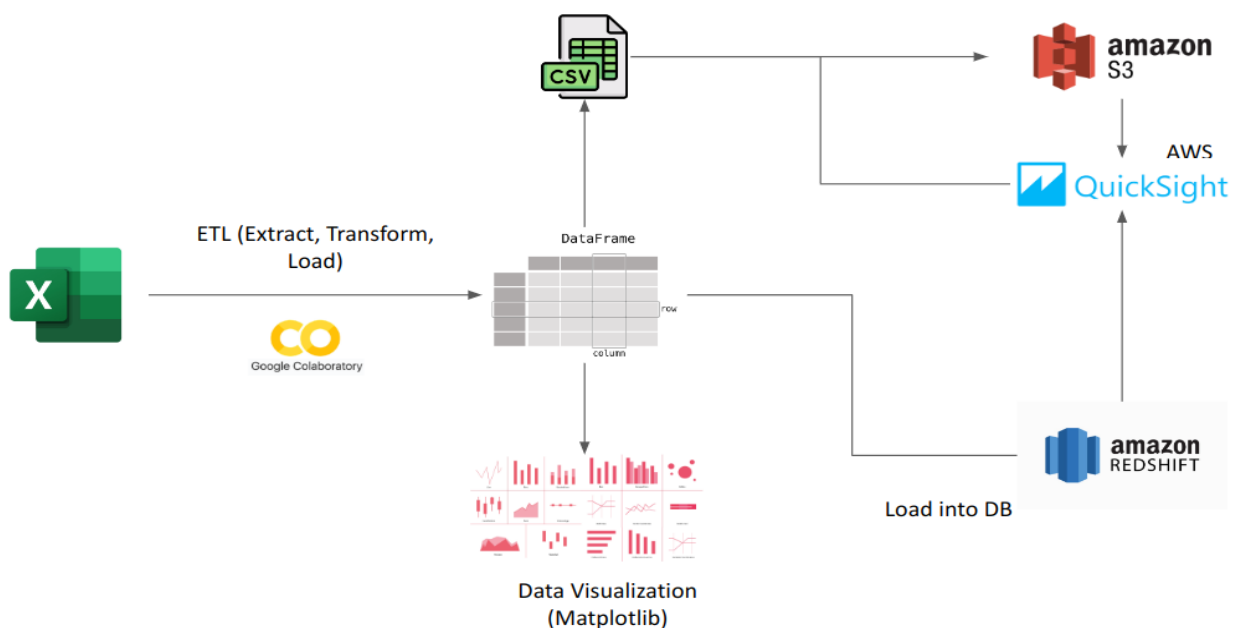


ETL Pipeline Documentation

The following ETL Pipeline is built using Python programming language and its associated tools & frameworks. The end-to-end pipeline works on cloud environments using Google Colab Notebook and Amazon Web Services (AWS) tools.

Below is the architecture of ETL pipeline being built -



Components Involved -

1. Google Colab Notebook (Python 3.9 configured)
2. AWS SDK (with associated AWS account)
 - a. AWS S3
 - b. Amazon Redshift
 - c. AWS QuickSight
3. Python packages
(Pandas, NumPy, Matplotlib, SQLAlchemy, boto3, psycopg2)


The steps of processes involved in creating and triggering the pipeline are as follows -

1. Setting up AWS components

Create AWS S3 Bucket

S3 Bucket will store the processed CSV files which can then be used as a source for Data Analytics and Visualization.

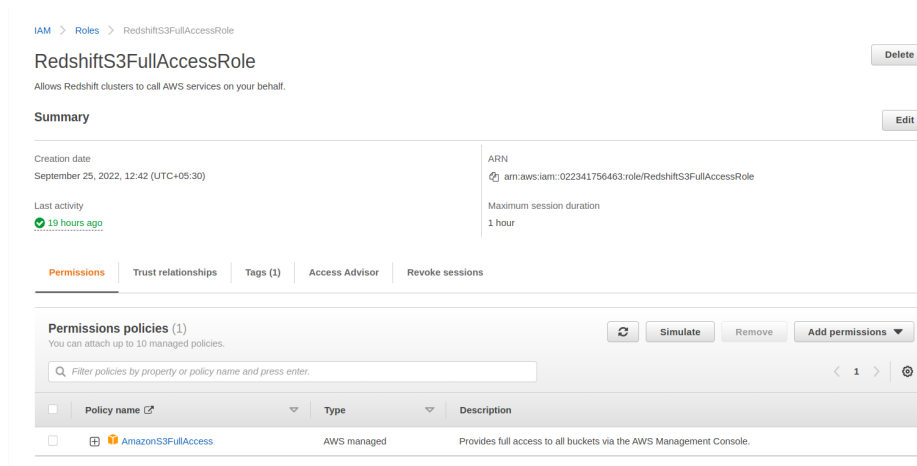
1. Create S3 bucket (public) using AWS Management Console
2. Using Python AWS SDK (boto3) connect with S3 API.
3. Bucket Name - **inputetlbucket**



Create AWS IAM Role

The IAM Role created will give access to Redshift to access S3 bucket files to load data if needed. (*optional)

1. IAM Role Name - **RedshiftS3FullAccessRole**
2. Policy Attached - **AmazonS3FullAccess**



2. Installing Python Packages

Since the ETL Pipeline is run using Google Colab, external python packages are essential.

The list of packages which are installed first are -

```
[ ] import boto3
import pandas as pd
import numpy as np
import psycopg2
import json
import matplotlib.pyplot as plt
import sqlalchemy as sa
from sqlalchemy.engine.url import URL
from sqlalchemy import orm as sa_orm
```

3. Python- AWS SDK Configuration

1. Using boto3 we will connect with EC2 resource, S3 resource, Redshift client and IAM resource for code integration in AWS platform.
2. EC2 resource is used to get the default VPC and attach that to AWS Redshift Cluster.
3. S3 client will be used to fetch bucket and add files to bucket during Load process of ETL pipeline
4. The Redshift client will be used to load data from dataframe to database.
5. IAM role will be used to perform the following actions -
 - a. Attach RedshiftS3FullAccessRole to Redshift
 - b. Get Redshift Cluster ARN to perform actions on it.

4. Create Redshift Cluster

Redshift Cluster is our target DB (Data Warehouse) to store the final dataset entity for Data Analytics and Visualization.

The created cluster is assigned into VPC configured default.

Cluster overview (1)		Any status ▾
Cluster	Status	
etl-cluster	Available	
		View all clusters

```
# Create Redshift Cluster
try:
    response = redshift.create_cluster(
        ClusterType=CLUSTER_TYPE,
        NodeType=NODE_TYPE,
        DBName=DB_NAME,
        ClusterIdentifier=CLUSTER_IDENTIFIER,
        MasterUsername=DB_USER,
        MasterUserPassword=DB_PASSWORD,
        IamRoles=[redshift_role_arn]
    )
except Exception as e:
    print(e)
```

5. Extract data from Spreadsheet

Pandas is used to extract the data from Spreadsheet (CSV or XLSX format) to create the data frame first. Data Nature is then described using pandas utility functions like `head()`, `isna()`, `mean()`, `duplicated()`, `info()`, `shape` etc.

Based on the nature of data and data types, it is then transformed from a raw data frame to a processed meaningful dataframe.

```
[ ] df = pd.read_excel('WS.xlsx')
```

```
▶ # data type che check for DB loading
print("Dataframe Shape : ", df.shape)
df.info()
```

```
↳ Dataframe Shape : (1629, 18)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1629 entries, 0 to 1628
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0    ID                                    1629 non-null   object
1    Company Name                         1629 non-null   object
2    Number of Employees                  1613 non-null   float64
3    % Developers Count                   1209 non-null   float64
4    Estimated Developers Count           1248 non-null   float64
5    Developers on LinkedIn               425 non-null    object
6    Business Description                 1608 non-null   object
7    Region of Headquarters               1611 non-null   object
8    Country of Headquarters              1611 non-null   object
9    Revenue                             1394 non-null   float64
10   Currency                            1629 non-null   object
11   TRBC Economic Sector Name            1606 non-null   object
12   TRBC Business Sector Name            1373 non-null   object
13   TRBC Industry Group Name             1371 non-null   object
14   URL                                  1591 non-null   object
15   Sales Person                         1629 non-null   object
16   Sales Person Email                   1629 non-null   object
17   Days Since First Contact              1620 non-null   float64
dtypes: float64(5), object(13)
memory usage: 229.2+ KB
```

```
▶ # check missing data
df.isna().mean()*100
```

```
↳ ID                                0.000000
Company Name                        0.000000
Number of Employees                  0.982198
% Developers Count                   25.782689
Estimated Developers Count           23.388582
Developers on LinkedIn               73.910374
Business Description                 1.289134
Region of Headquarters               1.104972
Country of Headquarters              1.104972
Revenue                             14.426028
Currency                             0.000000
TRBC Economic Sector Name            1.411909
TRBC Business Sector Name            15.715163
TRBC Industry Group Name             15.837937
URL                                  2.332719
Sales Person                         0.000000
Sales Person Email                   0.000000
Days Since First Contact              0.552486
dtype: float64
```

```
[ ] # check for duplicate data
df.duplicated().mean()*100
```

```
0.0
```

6. Transform Data

The set of processes performed during dataset transformation are -

1. Removing unwanted columns
2. Filling Empty cells (Nan, Null valued) with Aggregate values
3. Populating columns through peer dependent columns (A & B -> C)
4. Removing rows having empty cells (> 50-60%)
5. Bringing consistency and uniformity in column by maintaining a single data type.
6. Precising cell values using round functions and type casting.
7. Changing value of denormalized cells.

Some common utility functions used during the process includes -

`drop()`, `isna()`, `fillna()`, `loc()`, `value_counts()`,
`apply()`, `mean()`, `median()`, `astype()`

7. Load Data

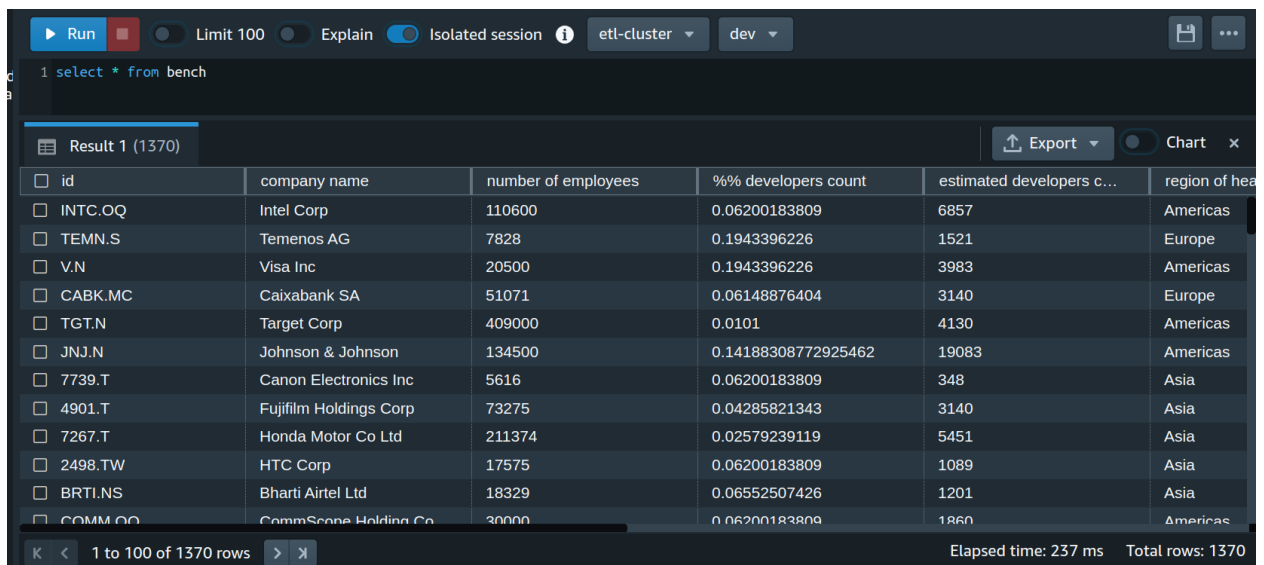
Transformed data is important in terms of insights and analytics so the data need to be stored in Database (DynamoDB) or Data Lake (S3) or Data Warehouse (Redshift).

1. Load data into Redshift

The rows and columns of dataframe are mapped with DB Schema and then rows are copied into redshift db.

- SQLAlchemy's utility tools like `create_engine()` and URL prepare to transfer data to DB. `to_sql()` commands copies data into the result set using the engine.
- Psycopg2 handles IO Errors during data loading.
- `Session()` is maintained for concurrent and multiple data loading use cases.

```
# load data to db
try:
    df.to_sql('bench', engine, index=False, if_exists='replace')
except psycopg2.Error as e:
    print('Error : Failed creating table !')
    print(e)
```



The screenshot shows a data visualization tool interface. At the top, there are controls for running a query, including a 'Run' button, a 'Limit 100' toggle, an 'Explain' toggle, an 'Isolated session' toggle, and dropdown menus for 'etl-cluster' and 'dev'. Below these controls, the SQL query '1 select * from bench' is entered. The results are displayed in a table with 6 columns: 'id', 'company name', 'number of employees', '% developers count', 'estimated developers c...', and 'region of head'. The table contains 1370 rows. The first few rows are visible, showing company names like Intel Corp, Temenos AG, Visa Inc, Caixabank SA, Target Corp, Johnson & Johnson, Canon Electronics Inc, Fujifilm Holdings Corp, Honda Motor Co Ltd, HTC Corp, Bharti Airtel Ltd, and CommScope Holding Co. The bottom of the interface shows a pagination bar with '1 to 100 of 1370 rows' and a status bar indicating 'Elapsed time: 237 ms' and 'Total rows: 1370'.

id	company name	number of employees	% developers count	estimated developers c...	region of head
INTC.OQ	Intel Corp	110600	0.06200183809	6857	Americas
TEMN.S	Temenos AG	7828	0.1943396226	1521	Europe
V.N	Visa Inc	20500	0.1943396226	3983	Americas
CABK.MC	Caixabank SA	51071	0.06148876404	3140	Europe
TGT.N	Target Corp	409000	0.0101	4130	Americas
JNJ.N	Johnson & Johnson	134500	0.14188308772925462	19083	Americas
7739.T	Canon Electronics Inc	5616	0.06200183809	348	Asia
4901.T	Fujifilm Holdings Corp	73275	0.04285821343	3140	Asia
7267.T	Honda Motor Co Ltd	211374	0.02579239119	5451	Asia
2498.TW	HTC Corp	17575	0.06200183809	1089	Asia
BRTI.NS	Bharti Airtel Ltd	18329	0.06552507426	1201	Asia
COMM.OQ	CommScope Holding Co	30000	0.06200183809	1860	Americas

Database schema

	Field	Type	NL	CMP
A	id	character varying(256)	NULL	lzo
A	company name	character varying(256)	NULL	lzo
#	number of employees	bigint	NULL	az64
#	%% developers count	double precision	NULL	none
#	estimated developers count	bigint	NULL	az64
A	region of headquarters	character varying(256)	NULL	lzo
A	country of headquarters	character varying(256)	NULL	lzo
A	trbc economic sector name	character varying(256)	NULL	lzo
A	trbc business sector name	character varying(256)	NULL	lzo
A	trbc industry group name	character varying(256)	NULL	lzo
#	revenue (usd)	double precision	NULL	none

2. Load data into S3

The processed data frame is also converted into a csv file which is then stored into AWS S3 bucket using boto3 module.

```
[ ] # save transformed file as csv and upload to AWS S3
df.to_csv("WS.csv",index=False)
s3.Bucket('inputetlbucket').upload_file('WS.csv','WS.CSV')
```

inputetlbucket [Info](#)

Objects

Properties

Permissions

Metrics

Management

Access Points

Objects (1)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 Inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Copy S3 URI

Copy URL

Download

Open

Delete

Actions

Create folder

Upload

Find objects by prefix

< 1 >

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	WS.CSV	CSV	September 26, 2022, 03:44:59 (UTC+05:30)	194.6 KB	Standard

8. Metrics and Visualization

The processed data frame as it is ready for data analytics , few metrics have been produced and those metrics and insights are visualized using Data visualization tools and libraries.

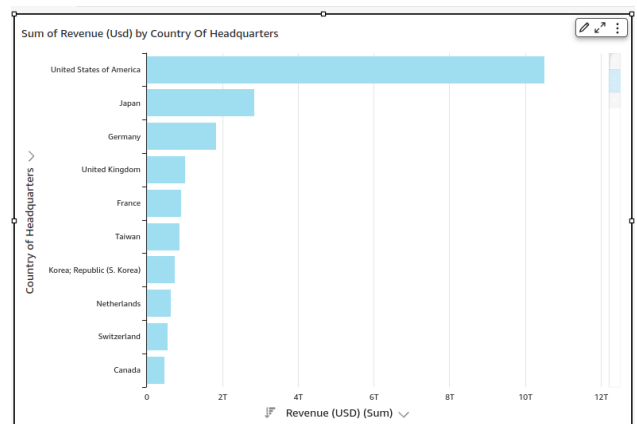

Some of the key metrics produced are as follows -


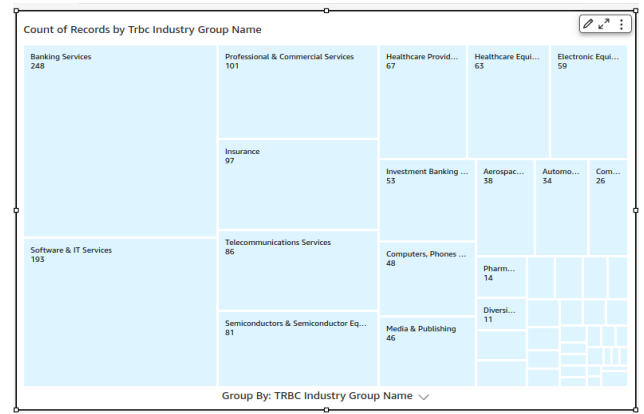

1. Total Revenue by Region of Headquarters
2. Total Revenue by Business Sector
3. Total Revenue by Technology Sector
4. Total Revenue by Economic Sector
5. Employes data by headquarter
6. Estimated developers count by Region of Headquarters
7. Sum of revenue by company

1. AWS QuickSight

AWS Quicksight is then finally connected to our ETL Pipeline which fetches data through various sources like Redshift, AWS S3 and CSV File.

The data is being fed to Quicksight which analyzes the data and creates data insights.





2. Matplotlib

Python libraries like Matplotlib are also used to visualize metrics derived from dataframe. Line Chart enhances data visualization showing trends of data.

